

Attention Is All You Need

Vaswani et al. 2017

Abstract

We propose a new network architecture, the Transformer, based solely on attention mechanisms, dispensing with

Experiments on two machine translation tasks show these models to be superior in quality while being more

1 Introduction

Recurrent neural networks have long been the de facto choice for sequence modeling.

The Transformer allows for significantly more parallelization and can reach a new state of the art in translation

2 Method

The model applies a multi-head self-attention mechanism.

Positional encodings are added to the input embeddings.